

# Continuous Deployment with Kubernetes basics

Wojciech Barczyński  
wbarczynski+trainings@gmail.com

## START

- minimum deployment, service, configmap, and secrets
- learn-as-you-go

# START

1. Structured logs
2. Liveness
3. Liveness and readiness
4. Prometheus metrics
5. Zero-downtime deployment

## 12FACTOR APPS

- Good to know what [12factorapp](#) is.
- 12factor is not new, so you might find points obvious.

# 12FACTOR APPS

Few that are very relevant for Kubernetes:

1. III. **Config** - environment variable, injected by environment
2. VIII. **Concurrency by process**
3. IX. **Disposability**

## 12FACTOR APPS

- 4. X. **Dev/Prod parity** including time, tooling, and backing services
- 5. XI. **Logs** - push your logs to stdout and let log routers to take care of them

## Continuous Deployment

- start at the
- see [Continuous Deployment slides](#)

# Probes

- liveness - am I dead?
- readiness - can I serve requests?

[slides on zero-downtime deployments](#)



## Probes

- deployment, statefulsets, daemonsets
- self-healing

# Liveness Probe

```
livenessProbe:
  httpGet:
    path: /model
    port: 8000
    httpHeaders:
      - name: X-Custom-Header
        value: Awesome
  initialDelaySeconds: 600
  periodSeconds: 5
  timeoutSeconds: 18
  successThreshold: 1
  failureThreshold: 3
```

[k8s docs](#)

## Liveness Probe

- pod gets restarted
- too many restarts = CrashLoop

[k8s docs](#)

## Liveness Probe

- I am dead
- My dependencies are dead → Readiness probe

## Liveness Probe

- Good to detect that we do not respond
- We can actively set it up (less common)

# Readiness Probe

```
readinessProbe:
  exec:
    command:
    - cat
    - /tmp/healthy
  initialDelaySeconds: 5
  periodSeconds: 5
```

[k8s docs](#)

## Readiness Probe

- temporary cannot server traffic

## Cannot Serve Traffic

- loading or reloading ML model
- stop getting requests before shutdown
- our deps are dead\*

[\*] after retries, usually the app kills itself.



## Probes best practices

- lightweight
- SHOULD NOT call external components

Pod lifecycle

**Zero-downtime deployments**

**Your app should on stop**

1. we get SIGTERM signal

## **Your app should on stop**

1. we get SIGTERM signal
2. app **MUST** make readinessProbe fail

## **Your app should on stop**

1. we get SIGTERM signal
2. app MUST make readinessProbe fail
3. k8s does not send more requests

## **Your app should on stop**

1. we get SIGTERM signal
2. app MUST make readinessProbe fail
3. k8s does not send more requests
4. app shuts down gracefully

## **Your app should on stop**

1. we get SIGTERM signal
2. app gives 500 on readinessProbe
3. k8s does not send more requests
4. app shuts down gracefully
5. kubernetes forces kill if 30s limit exceeded

# GRACEFUL SHUTDOWN

```
func (s *Service) prepareShutdown(h Server) {  
    signal.Notify(s.Stop, os.Interrupt, syscall.SIGTERM)  
    <-s.Stop  
    s.StatusNotReady()  
    shutdown(h)  
}
```

from [missy](#)



# Deployment Strategies

Demos and examples of the implementation:

zero-downtime deployments

# Observability

# Observability

- logging / sentry
- monitoring
- tracking

WHY?  
MONOLIT ;)





# WHY?

## MICROSERVICES ;)





# Observability

	Metrics	Logging	Tracing
CapEx	Medium	Low	High
OpEx	Low	High	Medium
Reaction	High	Medium	Low
Investigation	Low	Medium	High

Go for Industrial Programming by Peter Bourgon

# App Logs

Before we dive into observability.

# App Logs

- structured
- support setting LOG\_LEVEL[\*] per env

[\*] Deployment manifest.



# Logs - Golang

- logrus
- uber-go/zap

Pass logger in the context.

# Logs - Python

- structlog
- loguru

# Fieds

What do you think?

- url, retries, error code
- ?

# Centralized Logging

- ElasticSearch
- Fluentd
- Kibana

Worth following: [Loki](#)

# Centralized Logging

- Debugging tool
- Post-mortem
- Finding the needle
- High TCO

Sentry

# Monitoring

## SaaS / Cloud

METRIC

Name	Label	Value
traefik_requests_total	code="200", method="GET"	3001

from [slides](#)



# HOW TO FIND THE RIGHT METRIC?

- RED
- USE

## HOW TO FIND THE RIGHT METRIC?

- Measure from the client experience
- Alert on the client experience

## LOW LEVEL METRICS

- It is not onPremise :)
- Yes, if they correlate with experience [\*]

[\*] see saturation in [4 Golden signals](#)

# RED

---

Rate	How busy is your service?
------	---------------------------

---

Error	Errors
-------	--------

---

Duration	What is the latency of my service?
----------	------------------------------------

Tom Wilkie's guideline for instrumenting applications.

# RED

- **Rate** - how many request per seconds handled
- **Error**
- **Duration** (distribution)

# RED

- **Rate**
- **Error** - how many request per seconds handled we failed
- **Duration**

# RED

- **Rate**
- **Error**
- **Duration** - how long the requests took

# RED

- Follow Four Golden Signals by Google SREs [1]
- Focus on what matters for end-users

[1] Latency, Traffic, Errors, Saturation ([src](#))



# RED

Not recommended for:

- batch-oriented
- streaming services

# USE

---

Utilization	the average time that the resource was busy servicing work
-------------	--

---

Saturation	extra work which it can't service, often queued
------------	---

---

Errors	the count of error events
--------	---------------------------

---

Documented by [Berdan Gregg](#)

# USE

- Utilization: as a percent over a time interval:  
"one disk is running at 90% utilization".
- Saturation:
- Errors:

# USE

- **Utilization:**
- **Saturation:** as a queue length, e.g.,  
"the CPUs have an average run queue length of four".
- **Errors:**

# USE

- **utilization:**
- **saturation:**
- **errors:** scalar counts, e.g.,  
"this network interface drops packages".

# USE

- **traditionally** more instance oriented
- still useful in the microservices world

# Monitoring with prometheus

- Slides:  
[github.com/wojciech12/talk\\_monitoring\\_with\\_prometheus](https://github.com/wojciech12/talk_monitoring_with_prometheus)
- Prometheus naming conventions: [docs](#)

Questions?



# Survey

Your feedback is important:

- TBA